

CS5002 Project Write-Up

Yao Xiao

December 6, 2016

1 Overview

The program implements Huffman coding, a data compression algorithm that uses fixed-to-variable length encoding to exploit the non-uniform distribution of characters in files.

2 Changes

I implemented everything in the proposal, except the 'if time allows' part, because time was cruel and didn't allow.

In terms of implementation, I made a few changes - I didn't sort the frequency array, I used the data structure linked list which I didn't think of, and I modified the struct I used to represent the leaves. I also implemented the code in Dr. Racket with the help of SICP (Structure and Interpretation of Computer Programming).

3 User Interaction

On the command line, do the following:

```
> make  
> ./huffman
```

Follow the system prompt and proceed.

4 Results

I tried out the huffman tree with four files - test_small (421 bytes), test_even (712 bytes), test_mid (1492 bytes), and test_large (3408 bytes).

The compressed test_small file is of size 454 bytes, which is larger than the original file. This is because we also included the tree in the file. test_even stroke a balance - the compressed file is of size 711 bytes. This is not a robust analysis but I suspect that the break-even point is around 700 bytes. Then, when I go over 1KB, the advantage of Huffman coded files become more apparent. File test_mid achieved 21% decrease in size. File test_large achieved 30% decrease in size. The biggest file I tried, test_super (around 11 KB), achieved 36% decrease in size. So this is definitely working!

5 Lessons Learnt

META

1. In Theory vs. In Code

Huge difference! It took me about ten times the effort to implement than to understand the algorithm. It was a humbling experience. And it's a little scary - I knew how the algorithm should work in every step and only had to worry about using C to make it real, and I knew for certain that there is an answer, this is a very basic compression algorithm, and still it took me so long to succeed. I can only imagine what the process is like for real research - it's doing this over and over again, for problems that have no clear solution.

2. Code Validation

Just because the code compiles does not mean it does what you want it to do. Always write tests with synthetic data.

3. Exploit the strengths of a language

Exploit the data structures in different languages - use Racket for easy implementation to solidify understanding, use C for speed and low-level manipulation. Exploit the deterministic nature of Huffman tree when saving the tree.

4. Simplicity vs. Speed

Simple code sometimes has a high price tag (see point 3 for detail).

5. There is nothing more motivating than an impending deadline.

CODING

1. The real challenge of the project was to use the bare-bone data structures provided in C. The ubiquitous list structure in Racket is actually a linked list in disguise. I realized this when I was trying to build an array of pointers in C. I spent a lot of time writing (and debugging) basic functions such as finding the minimum element and removing an element from a linked list. I missed writing one-liners in Racket.

However, the one-liners in Racket turned out to be run-time landmines - for example, the `member?` function that checks if an element is in a list is actually $O(n)$ because it traverses the entire linked list. To avoid this in C, I took advantage of the speed of array indexing - I added an array of length 256 in the tree struct to track what elements is in a tree.

If I were to choose between C and Racket, I would choose C, because even though there are less readily available data structures made for me, I have a lot more control over the run time of the program.

2. This is the first time I worked with a real file in C. That was the most time-consuming part of the project. I found out about ASCII encoding vs. UNICODE. I discovered the usefulness of the seemingly useless bell character. I learnt hexdump. I learnt bit writing. I have come to appreciate simple, bare-bone text editors.
3. I have finally gotten a taste of the evils of `scanf()` - it leaves stuff in the buffer! So later if you want to get input again, you will be stuck with the junk left by `scanf()` unless you clean it. I could have done a better job by using `fgetc()` but I was lazy so I used a quick fix (thanks Stack Overflow!).
4. I got better at debugging. I got better at accepting the existence of bugs. I especially got better at locating segmentation faults. One thing I did not do during the project was aggressively checking for errors because I was lazy and only did it at the end as a touch-up. It could have saved me hours of debugging.